

# PROGRAMMING CONCEPTS AND LANGUAGES, 2024w

**Selected Topics: Mojo, Rust**

# SELECTED TOPICS

## **Mojo Programming Language**

# CONTENTS

## The Mojo Programming Language

### Basics

- Hello World
- Functions, values, variables, etc.

### Mojo Ownership System

### Literature

- <https://docs.modular.com/mojo/manual/basics>
- <https://docs.modular.com/mojo/playground>

# MOJO PROGRAMMING LANGUAGE

## High-performance systems programming language

- Open-source standard library, may become fully open-source as it matures\*\*
- Superset of Python
  - Makes Mojo easier to learn, while providing compatibility with Python libraries
- Strong type checking (optional)
- Memory safety
- Aiming for performance that is on par with C++ and CUDA without the complexity
- Built-in utilities for optimization, parallelization, auto-tuning

## Still in development and many features are missing

- Modular\* team builds on their experience building other technologies like Clang and Swift programming language
- Focus on compilation model and systems programming features (not necessarily syntax)
- Also uses lessons learned from languages like Rust, Swift, Julia, Zig, Nim, etc.

\*Modular is a company founded by Chris Lattner (Swift, LLVM)

- See <https://www.modular.com>
- See <https://www.modular.com/blog/the-next-big-step-in-mojo-open-source>

# MOJO :: PERFORMANCE SHOWCASE

LANGUAGES	TIME (S)*	SPEEDUP VS PYTHON
Python 3.10.9	1027s	1X
PYPY	46.1s	22x
Scalar C++	0.20s	5,000x
Mojo 	0.03s	68,000x

Algorithm: **Mandelbrot**

Instance: **AWS C1.xlarge Intel Xeon**

## Sources

- <https://www.modular.com/mojo>
- <https://www.modular.com/blog/mojo-a-journey-to-68-000x-speedup-over-python-part-3>

# MOJO :: GOALS/HIGHLIGHTS

## Full compatibility with the Python ecosystem

- Python is dominating in ML (TensorFlow, PyTorch) and Mojo also aims to address AI/ML challenges
  - Python is very productive, but anything that needs to go fast is written in C/C++

## Predictable low-level performance and low-level control

## Targeting heterogeneous hardware with a single language

- Ability to deploy subsets of code to accelerators
- Addressing AI/ML challenges
  - AI/ML very often used on accelerators (e.g., GPUs)
  - CPUs are equally important, not only for operations accelerators cannot handle
- Code needs to be optimized at compile time for various processing units (CPUs, GPUs, and other accelerators), which may all support different types of operations

## Utilize the next-generation compiler technologies

- Built from the ground-up using **Multi-level Intermediate Representation (MLIR\*)**
  - A compiler infrastructure that's ideal for heterogeneous hardware, from CPUs and GPUs, to various AI ASIC

\*we will come back to MLIR soon

**some background before we move on**

# PYTHON, CPYTHON

## Python

- High-level, interpreted programming language
- Python language specification – a document that describes the Python language

## CPython

- Reference implementation of Python, written in C
- Interprets Python code and compiles it into bytecode, which is executed by the CPython virtual machine
- Supports a wide range of libraries and is the most compatible with Python's ecosystem
- The compiler, standard library modules, core types, test suite, shared language specification, ...

## Pros

- A dominant force in AI/ML and many other fields, easy to learn, community support, packages, tooling, can be used as a “glue layer” and allows building libraries in C/C++, etc.

## Cons

- Poor low-level performance, single-threaded due to global interpreter lock (GIL) in Cpython\*
- Using C/C++ for performance while building libraries requires a good knowledge of internal Python and C/C++ (and perhaps CUDA and other technologies)

\*There have been some improvements in this field recently



# MOJO AND PYTHON

## Full compatibility with **Python** and existing libraries

- However, it is not there yet!

## Uses **CPython** only for interoperability

- Python 3 code without modification with full compatibility with the ecosystem (for migration and adoption purposes)
- Pure Mojo does not rely on any other runtime or compiler technology!

## Mojo aims to address the need for a **unified language** that eliminates the need to rely on C/C++ within Python libraries

- Mojo aims for highest performance possible
- Can target different accelerators with predictability and control over how computation happens
- Uses an **MLIR**-based infrastructure to enable high-performance execution on a wide range of hardware

# LLVM\*

**Started as a research project at the University of Illinois**

**Aim to support both static and dynamic compilation of arbitrary programming languages**

## **LLVM Project**

- A collection of modular and reusable, industrial strength, compiler and toolchain technologies
- Code in the LLVM project is licensed under the "Apache 2.0 License with LLVM exceptions"

## **LLVM Core**

- Libraries that provide a modern source- and target-independent optimizer, and code generation support for many popular CPUs
- Build around LLVM intermediate representation ("**LLVM IR**")

## **Clang (a frontend)**

- A high-end C/C++/Objective C compiler

**LLDB, libc++, compiler-rt, **MLIR**, OpenMP, ...**

\*Originally stood for Low Level Virtual Machine, now the name is no longer an initialism (so just LLVM)

# LLVM :: SOME INFO ON HOW COMPILERS WORK

## (Clang) Frontend

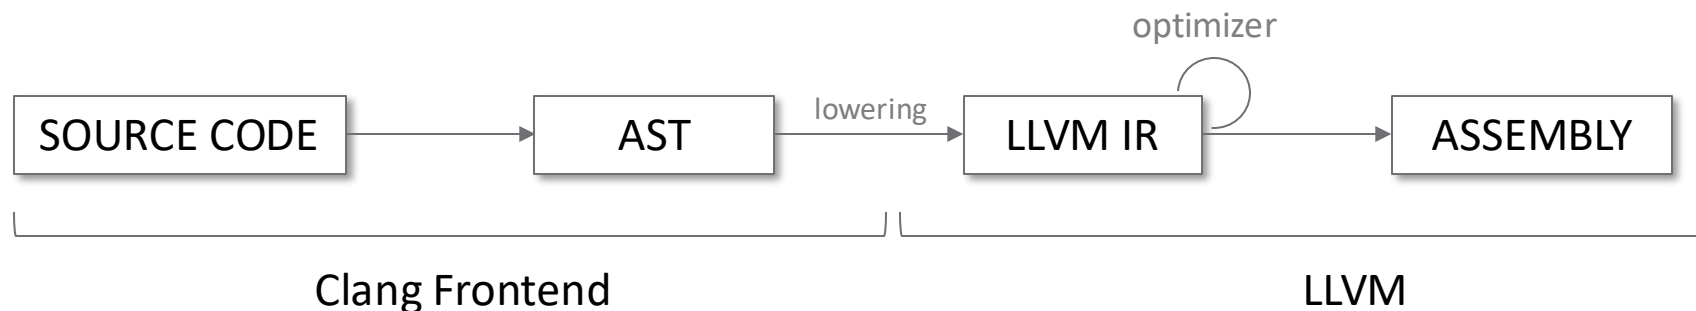
- Parses source code into an Abstract Syntax Tree (AST) and lowers it into an intermediate representation (IR)

## LLVM Optimizer

- An optimizer (or “middle-end”) that transforms IR into better IR

## LLVM Backend

- Converts IR into machine code for a particular platform



# LLVM :: RUST EXAMPLE :: SOURCE CODE

```
pub fn square(x: i32) -> i32 {  
    x * x  
}
```

# LLVM :: RUST EXAMPLE :: IR

```

@alloc_9be5c135c0f7c91e35e471f025924b11 = private unnamed_addr constant <{ [15 x i8] }> <{ [15 x i8] c"/app/example.rs" }>, align 1
@alloc_1d63fdc7829b250028d05847a3ed2432 = private unnamed_addr constant <{ ptr, [16 x i8] }> <{ ptr
@alloc_9be5c135c0f7c91e35e471f025924b11, [16 x i8] c"\0F\00\00\00\00\00\00\02\00\00\00\03\00\00\00" }>, align 8

define i32 @example::square::h658e3574c1514f63(i32 %x) unnamed_addr #0 !dbg !7 {
start:
%0 = call { i32, i1 } @llvm.smul.with.overflow.i32(i32 %x, i32 %x), !dbg !12
%_2.0 = extractvalue { i32, i1 } %0, 0, !dbg !12
%_2.1 = extractvalue { i32, i1 } %0, 1, !dbg !12
br i1 %_2.1, label %panic, label %bb1, !dbg !12

bb1:
ret i32 %_2.0, !dbg !13

panic:
call void @core::panicking::panic_const::panic_const_mul_overflow::h446c6be747f81624(ptr align 8 @alloc_1d63fdc7829b250028d05847a3ed2432)
#3, !dbg !12
unreachable, !dbg !12
}

declare { i32, i1 } @llvm.smul.with.overflow.i32(i32, i32) #1

declare void @core::panicking::panic_const::panic_const_mul_overflow::h446c6be747f81624(ptr align 8) unnamed_addr #2

attributes #0 = { nonlazybind uwtable "probe-stack"="inline-asm" "target-cpu"="x86-64" }
attributes #1 = { nocallback nofree nosync nounwind speculatable willreturn memory(none) }
attributes #2 = { cold noline noreturn nonlazybind uwtable "probe-stack"="inline-asm" "target-cpu"="x86-64" }
attributes #3 = { noreturn }

```

Src: <https://godbolt.org>

# MOJO :: MULTI-LEVEL INTERMEDIATE REPRESENTATION (MLIR)

**Compiler framework (started at Google, now LLVM Developers)**

**Hybrid IR to support “multiple different requirements in a unified infrastructure”**

- The ability to represent **dataflow graphs** (such as in TensorFlow), including dynamic shapes, the user-extensible op ecosystem, TensorFlow variables, etc.
- Optimizations and transformations typically done on such **graphs** (e.g., in Grappler)
- Ability to host **high-performance-computing-style** loop optimizations across kernels (fusion, loop interchange, tiling, etc.), and to **transform memory layouts of data**.
- Code generation “lowering” transformations such as DMA insertion, explicit cache management, memory tiling, and vectorization for 1D and 2D register architectures
- Ability to represent target-specific operations, e.g., accelerator-specific high-level operations
- Quantization and other graph transformations done on a Deep-Learning graph
- Polyhedral primitives
- Hardware Synthesis Tools / HLS

# MOJO :: HELLO WORLD

Mojo requires a `main()` function declared with `fn` or `def`.

```
def main():  
    # CHECK: Hello Mojo 🔥!  
    print("Hello Mojo 🔥!")  
    for x in range(9, 0, -3):  
        print(x)
```

Output:

```
Hello Mojo 🔥!  
9  
6  
3
```

How to run on your machine?

- <https://docs.modular.com/mojo/manual/get-started>

# MOJO :: BLOCKS, STATEMENTS, COMMENTS

## Blocks and statements

- Code blocks such as functions, conditions, and loops are defined with a colon followed by indented lines

```
def loop():  
    for x in range(5):  
        if x % 2 == 0:  
            print(x)
```

## Comments

- One-line comments with the hash # symbol
- Multi-line comments with triple quotes """ This is a comment """

```
# this is a comment  
  
"""  
This is a  
comment  
"""
```



# MOJO :: FUNCTIONS

## Two styles with minor differences

- Both styles support optional, keywords and variadic arguments

### fn (Rust-style)

- Enforces type-checking and memory-safe behaviors
- Compile-time checks to ensure the function receives and returns the correct types
- Arguments passed by immutable reference (default)

```
fn greet2(name: String) -> String:  
    return "Hello, " + name + "!"
```

### def (Python-style)

- Allows no type declarations and dynamic behaviors
- Arguments passed by value

```
def greet(name):  
    return "Hello, " + name + "!"
```

# MOJO :: FUNCTIONS :: FN

## fn (Rust-style)

- Arguments **must specify a type**
  - except for the self argument in struct methods
- Return values **must specify a type**, unless the function doesn't return a value
- If you don't specify a return type, it defaults to **None** (meaning no return value)
- By default, arguments are received as an immutable reference
  - values are read-only, using the borrowed argument convention
  - This prevents accidental mutations, and permits the use of non-copyable types as arguments
- If you want a local copy, you can simply assign the value to a local variable.
  - You can also get a mutable reference to the value by declaring the **inout argument convention**
- If the function raises an exception, it must be explicitly declared with the raises keyword
  - A def function does not need to declare exceptions

# MOJO :: FUNCTIONS :: DEF

## def (Python-style)

- Like in Python, but you can choose if you want to specify argument and return types
- Arguments don't require a declared type
  - Undeclared arguments are actually passed as an object, which allows the function to receive any type
    - Mojo infers the type at runtime
- Return types don't need to be declared and also default to **object\***
  - If a def function doesn't declare a return type of **None**, it's considered to return an **object by default**
  - If an argument is an object type, it's received as a reference, following object reference semantics
  - If an argument is any other declared type, it's received as a value

\*The object type allows for dynamic typing because it can represent any type in the Mojo standard library, and the actual type is inferred at runtime.

```
def greet(name):  
    greeting = "Hello, " + name + "!"  
    return greeting
```

def (Python-like)

```
def greet(name: String) -> String:  
    var greeting = "Hello, " + name + "!"  
    return greeting
```

def (with types)

# MOJO :: VARIABLES

## A name that holds a value or object

- Mutable by default
- Constant values can be defined with the alias keyword
- Support type annotations
- Late Initialization
- Implicit type conversion

## Explicitly or implicitly declared

- Type is inferred (strongly typed)
  - You cannot assign a value of a different type
- Some types support implicit conversion
  - e.g., int to float
- Examples:

```
var user_id: Int = "Sam" #does not compile
count = 8 # count is type Int
count = "Nine?" # Error: can't implicitly convert 'StringLiteral' to 'Int'
```

```
var a = 5
var b: Float64 = 3.14
```

Explicitly declared

```
a = 5
b = 3.14
```

Implicitly declared

```
var temperature: Float64 = 99
print(temperature)
```

99.0

# MOJO :: VARIABLE SCOPE

## Explicitly declared variables

- Lexical scoping
- Variables declared with `var` are bound by lexical scoping
  - Nested code blocks can read and modify variables defined in an outer scope
  - An outer scope cannot read variables defined in an inner scope at all

```
def lexical_scopes():
    var num = 1
    var dig = 1
    if num == 1:
        print("num:", num) # Reads the outer-scope "num"
        var num = 2 # Creates new inner-scope "num"
        print("num:", num) # Reads the inner-scope "num"
        dig = 2 # Updates the outer-scope "dig"
    print("num:", num) # Reads the outer-scope "num"
    print("dig:", dig) # Reads the outer-scope "dig"
```

```
num: 1
num: 2
num: 1
dig: 2
```

## Implicitly declared variables

- Scoped at function level
- If the value of an implicitly-declared variable inside the `if` block is changed, it actually changes the value for the entire function

```
def function_scopes():
    num = 1
    if num == 1:
        print(num) # Reads the function-scope "num"
        num = 2 # Updates the function-scope variable
        print(num) # Reads the function-scope "num"
    print(num) # Reads the function-scope "num"
```

```
1
2
2
```

# MOJO :: TYPES, NUMERIC TYPES

## All values in Mojo have an associated data type

- Most of the types are nominal (or “named”) types, defined by a struct
- Some types that are not defined as structs
  - Functions are typed based on their signatures.
  - **NoneType** is a type with one instance, the **None** object, which is used to signal “no value”
- Basic types are included, e.g., numeric values, strings, boolean values and others
- The standard library also includes many more types that you can import as needed
  - E.g., collection types, utilities for interacting with the filesystem and getting system information, etc.

## Booleans

- Either **True** or **False**
- Negated with **not**

## Numeric Types

- Int8, UInt8, Int16, UInt16, Int32, UInt32, Int64, UInt64, Float16, Float32, Float64
- All support number and bitwise operators, usable with the **math** module
- Numeric literals IntLiteral and FloatLiteral, which are used at compile time

# MOJO :: SIMD AND DTYPE

## **SIMD type as a basis for numeric types (SIMD – single instruction, multiple data)**

- A processor technology that allows performing an operation on an entire set of operands at once
- Fixed-size array of values that can fit into processor's register
- Operations on SIMD values are applied elementwise, on each individual element in the vector
- A DType value defining the data type in the vector (for example, 32-bit floating-point numbers)

```
var vec1 = SIMD[DType.int8, 4](2, 3, 5, 7)
var vec2 = SIMD[DType.int8, 4](1, 2, 3, 4)

var product = vec1 * vec2

print(product)
```

```
[2, 6, 15, 28]
```

## **In Mojo, a scalar type is a SIMD vector with a single element**

- Language design decisions, essentially all operations scalar or SIMD go through exactly the same code path

# MOJO :: STRINGS

## String type represents a mutable string

- Different from Python's standard string, which is immutable
- Support a variety of operators and common methods
- **StringLiteral** represents literal strings in the program source
  - Needs to be manually converted to String with `str()` method

```
var s: String = "Testing"  
s += " Mojo strings"  
print(s)
```

```
Testing Mojo strings
```

## “Stringable” trait

- Most standard type conform to it
- Represents a type that can be converted to a string
- **str(value)** can be used for explicit conversion

```
var s = str("Items in list: ") + str(5)  
print(s)
```



# MOJO :: COLLECTIONS :: LIST

## A dynamically-sized array of items

- Elements must be copyable and movable (**CollectionElement** trait)
- Element type can be passed as a parameter (`var l = List[String]()`)
- You can add, pop and access list elements

```
from collections import List

var list = List(2, 3, 5)
list.append(7)
list.append(11)
print("Popping last item: ",list.pop())
for idx in range(len(list)):
    print(list[idx], end=", ")
```

```
Popping last item: 11 2, 3, 5, 7,
```

# MOJO :: COLLECTIONS :: DICT

## An associative array of key-value pairs

- Key type and Value type specified as parameters
- The key type must conform to **KeyElement** trait
- The value elements must conform to **CollectionElement** trait
- **Dict** iterators all yield references, so you need to use the dereference operator []

```
from collections import Dict

var d = Dict[String, Float64]()
d["plasticity"] = 3.1
d["elasticity"] = 1.3
d["electricity"] = 9.7
for item in d.items():
  print(item[].key, item[].value)
```

```
plasticity 3.1000000000000001
elasticity 1.3
electricity 9.6999999999999993
```

# MOJO :: COLLECTIONS :: SET

## An unordered collection of unique items

- Must conform to the **KeyElement** trait.
- Adding, removing elements and testing whether a value exists in the set
- Unions and intersections between two sets

```
from collections import Set

i_like = Set("sushi", "ice cream", "tacos", "pho")
you_like = Set("burgers", "tacos", "salad", "ice cream")
we_like = i_like.intersection(you_like)

print("We both like:")
for item in we_like:
    print("-", item[])
```

```
We both like:
- ice cream
- tacos
```

# MOJO :: COLLECTIONS :: OPTIONAL

## Represents a value that may or may not be present

- Evaluates to True when it holds a value, False otherwise
- If it holds a value it can be retrieved using **value()** method, but calling it on an **Optional** with no value results in an undefined behavior
  - Similar to Rust, there is an **or\_else** method which returns the stored value if there is one, or a user-specified default value

```
var custom_greeting: Optional[String] = None
print(custom_greeting.or_else("Hello"))

custom_greeting = str("Hi")
print(custom_greeting.or_else("Hello"))
```

```
Hello
Hi
```

# MOJO :: CONTROL FLOW :: IF

## The `if` statement

- Short-circuit evaluation semantics
  - If the first argument to an `or` operator evaluates to `True`, the second argument is not evaluated
  - If the first argument to an `and` operator evaluates to `False`, the second argument is not evaluated

```
temp_celsius = 25
if temp_celsius > 20:
    print("It is warm.")
    print("The temperature is", temp_celsius * 9 / 5 + 32,
          "Fahrenheit." )
```

```
It is warm.
The temperature is 77.0 Fahrenheit.
```

## Conditional expressions

```
temp_celsius = 15
forecast = "warm" if temp_celsius > 20 else "cool"
print("The forecast for today is", forecast)
```

```
The forecast for today is cool
```

# MOJO :: CONTROL FLOW :: WHILE

## The **while** statement

- Repeatedly executes a code block while a given boolean expression evaluates to True
- A **continue** statement skips execution of the rest of the code block and resumes with the loop test expression
- A **break** statement terminates execution of the loop
- Can include else statement, which executes when loop's Boolean condition evaluates to False

```
n = 0
while n < 5:
    n += 1
    if n == 3:
        continue
    print(n, end=", ")
```

1, 2, 4, 5,

```
n = 0
while n < 5:
    n += 1
    if n == 3:
        break
    print(n, end=", ")
```

1, 2,

```
n = 5
while n < 4:
    print(n)
    n += 1
else:
    print("Loop completed")
```

Loop completed

# MOJO :: CONTROL FLOW :: FOR

## The **for** statement

- Iterates over a sequence, executing a code block for each element in the sequence
- Can iterate over any type that implements an `__iter__()` method that returns a type that defines `__next__()` and `__len__()` methods
- Can use **continue** and **break** statements
- Can use **else** clause

## For example, iterating over a Mojo List

```
from collections import List

states = List[String]("California", "Hawaii", "Oregon")
for state in states:
    print(state[])
```

```
California
Hawaii
Oregon
```

# MOJO :: STRUCTS

A data structure that allows you to encapsulate *fields* and *methods* that operate on an abstraction, such as a data type or an object

- Similar to a class in Python, but completely static (bound at compile-time – no runtime changes)
- *Fields* are variables that hold data relevant to the struct
- *Methods* are functions inside a struct that generally act upon the field data.

```
struct MyPair:
    var first: Int
    var second: Int

fn __init__(inout self, first: Int, second: Int):
    self.first = first
    self.second = second
```

```
var mine = MyPair(2,4)
print(mine.first)
```

Once you have a constructor, you can create an instance of MyPair and set the fields

## Can have methods and static methods

- A static method can be called without creating an instance of the struct
- Does not receive the implicit “self” argument, so it cannot access any fields of the struct
  - With @staticmethod decorator
- Special methods (or dunder methods)
  - Predetermined for special tasks (e.g., \_\_init\_\_(), \_\_del\_\_(), \_\_copyinit\_\_(), \_\_moveinit\_\_())



# MOJO :: STRUCTS VS PYTHON CLASSES

## **Python classes are dynamic**

- Allow for dynamic dispatch, monkey-patching (or “swizzling”), and dynamically binding instance fields at runtime

## **Mojo structs are static**

- Structures and contents of a struct are set at compile time and can't be changed while the program is running → you cannot add methods at runtime
- Structs allow you to trade flexibility for performance while being safe and easy to use
- The program knows exactly where to find the structs information and how to use it without any extra steps or delays at runtime

## **Mojo structs do not support inheritance ("sub-classing"), but a struct can implement traits.**

- Python classes support class attributes - values that are shared by all instances of the class, equivalent to class variables or static data members in other languages
- Mojo structs don't support static data members

# MOJO :: TRAITS

## A trait is a set of requirements that a type must implement

- Similar to a template of characteristics for a struct
- If you want a struct with the characteristics defined in a trait, you must implement each characteristics (e.g., each method)
- Similar to Java interfaces, C++ concepts, Swift protocols, and Rust traits

## You can think of it as a contract

- A type that conforms to a trait guarantees that it implements all of the features of the trait

```
trait SomeTrait:  
    fn required_method(self, x: Int): ...
```

A trait

```
@value  
struct SomeStruct(SomeTrait):  
    fn required_method(self, x: Int):  
        print("hello traits", x)
```

A struct that conforms to the trait above

```
fn fun_with_traits[T: SomeTrait](x: T):  
    x.required_method(42)  
  
fn use_trait_function():  
    var thing = SomeStruct()  
    fun_with_traits(thing)
```

A function that uses the trait as an argument type

# MOJO :: PARAMETRIZATION

**A parameter is a compile-time variable that becomes a runtime constant**

- For performance reasons
- Declared in square brackets on a function or struct
- A function argument is a runtime value that is declared in parenthesis

```
fn repeat[count: Int](msg: String):  
    @parameter  
    for i in range(count):  
        print(msg)
```

**This function has**

- One parameter of type **Int**
- One argument of type **String**

**To call the function, you need to specify both the parameter and the argument**

```
fn call_repeat():  
    repeat[3]("Hello")  
    # Prints "Hello" 3 times
```

# Mojo: Value Semantics, Ownership

# MOJO :: VALUE SEMANTIC

## Mojo doesn't enforce value semantics or reference semantics.

- Allows each type to define how it is created, copied, and moved (if at all)
- You can choose between copy and reference semantics
- Default: value semantics, with tight control over reference semantics to avoid memory errors
- Note: Python is not value semantics

**Each variable has unique access to a value, and any code outside the scope of that variable cannot modify its value**

**Numeric values in Mojo are value semantic because they are trivial types, which are cheap to copy.**

```
x = 1
y = x
y += 1

print(x)
print(y)
```

```
1
2
```

# MOJO :: OWNERSHIP

**Set of rules that govern memory management**

## **The fundamental ownership rules**

- Every value has only one owner at a time
- When the lifetime of the owner ends, Mojo destroys the value
- If there are outstanding references to a value, Mojo keeps the value alive

**In the future, the Mojo lifetime checker will enforce reference exclusivity, so that only one mutable reference to a value can exist at a time.**

- This is not currently enforced (remember Rust?)

# MOJO :: OWNERSHIP

## **Prevents memory safety issues such as**

- Trying to free memory that has already been freed (double-free)
- Memory leaks (not freeing memory that should be freed)
- Use-after-free (accidentally deallocate data before the program is done with it)

## **Makes sure that there is only one “owner” for a given value at the time**

- When a value’s lifetime ends, Mojo calls its destructor, which is responsible for deallocating
- When the life span of the owner ends, Mojo destroys the value
- Programmers are still responsible for making sure any type that allocates resources (including memory) also deallocates those resources in its destructor. Mojo's ownership system ensures that destructors are called promptly

## **Lower performance overhead**

- Compared to automatic memory management (e.g., “garbage collector”)
- Relieves programmers of the burden of manual memory management and associated errors

# MOJO :: LIFETIMES

## Lifetimes

- Meaning: span of time when the value is valid
- Also refers to a specific type of parameter value used to help track the lifetimes of values and references to values

## Lifetime checker

- a compiler pass that analyzes dataflow through your program and identifies when variables are valid and inserts destructor calls when a value's lifetime ends

**Mojo uses lifetime values to track the validity of references. Specifically, a lifetime value answers two questions:**

- What logical storage location "owns" this value?
- Can the value be mutated using this reference?

**Most of the time, lifetime values are handled automatically by the compiler, but In some cases direct interaction with lifetime values is needed**

- When working with references - specifically ref arguments and ref return values
- When working with types like Reference or Span which are parameterized on the lifetime of the data they refer to



# MOJO :: OWNERSHIP

## **Mojo aims to provide full value semantics by default**

- Code quality and performance is dependent upon how functions treat argument values
- Value semantics provides consistent and predictable behavior, but reference semantics is required for a systems programming language for full control over memory optimizations
- Reference semantics used in a way that ensures all code is memory safe by tracking the lifetime of every value and destroying each one at the right time (and only once)
- Achieved through the use of argument conventions that ensure every value has only one owner at a time

## **Argument convention**

- Specifies whether an argument is mutable or immutable
- Whether the function owns the value

# MOJO :: OWNERSHIP :: ARGUMENT CONVENTION

**Defined by a keyword at the beginning of an argument declaration:**

- **borrowed**: The function receives an immutable reference
  - The function can read the original value (it is not a copy), but it cannot mutate (modify) it. `def` functions treat this differently, as described below.
- **inout**: The function receives a mutable reference
  - The function can read and mutate the original value (it is not a copy)
- **owned**: The function takes ownership
  - This means the function has exclusive ownership of the argument
  - Often, this also implies that the caller should transfer ownership to this function, but that's not always what happens and this might instead be a copy (as you'll learn below)
- **ref**: The function gets a reference with an associated lifetime
  - The reference can be either mutable or immutable
  - You can think of `ref` arguments as a generalization of the `borrowed` and `inout` conventions

```
fn add(inout x: Int, borrowed y: Int):  
    x += y  
  
fn main():  
    var a = 1  
    var b = 2  
    add(a, b)  
    print(a) # Prints 3
```

# MOJO :: SUMMARY

## Systems programming language

**Still in early development, some features are missing, for example:**

- Private/public
- Extended support for traits
- Classes
- C/C++ interoperability
  - Integration to transparently import Clang C/C++ modules. Mojo's type system and C++'s are very compatible
- Calling Mojo from Python
- Full MLIR decorator reflection
- No list or dict comprehensions
  - Mojo does not yet support Python list or dictionary comprehension expressions, like `[x for x in range(10)]`.
- No lambda syntax
- No async for or async with
  - Although Mojo has support for async functions with `async fn` and `async def`, Mojo does not yet support the `async for` and `async with` statements.
- Parameter closure captures are unsafe references
- Nested functions cannot be recursive
- ... (see <https://docs.modular.com/mojo/roadmap>)

# MOJO :: EXAMPLES

**Note that Mojo provides a standard library**

- <https://docs.modular.com/mojo/lib>

**Some interesting functions**

- <https://docs.modular.com/mojo/stdlib/algorithm/functional/parallelize>
- <https://docs.modular.com/mojo/stdlib/algorithm/functional/vectorize>
- <https://docs.modular.com/mojo/stdlib/algorithm/functional/stencil>
- <https://docs.modular.com/mojo/stdlib/algorithm/functional/tile>

**We follow the following links:**

- <https://docs.modular.com/mojo/notebooks/Mandelbrot>
- <https://www.modular.com/blog/how-mojo-gets-a-35-000x-speedup-over-python-part-1>
- <https://www.modular.com/blog/how-mojo-gets-a-35-000x-speedup-over-python-part-2>
- <https://www.modular.com/blog/mojo-a-journey-to-68-000x-speedup-over-python-part-3>

# MOJO :: RESOURCES, REFERENCES AND LINKS

## Resources at modular.com

1. <https://www.modular.com/mojo>
2. <https://docs.modular.com/mojo/manual/get-started>
3. <https://docs.modular.com/mojo/playground>
4. <https://docs.modular.com/mojo/lib>
5. <https://docs.modular.com/mojo/notebooks/>
6. <https://docs.modular.com/mojo/manual/>

# SELECTED TOPICS

## **Rust: Actix-Web**

# RUST :: ACTIX WEB

## An HTTP server and a web framework written in Rust

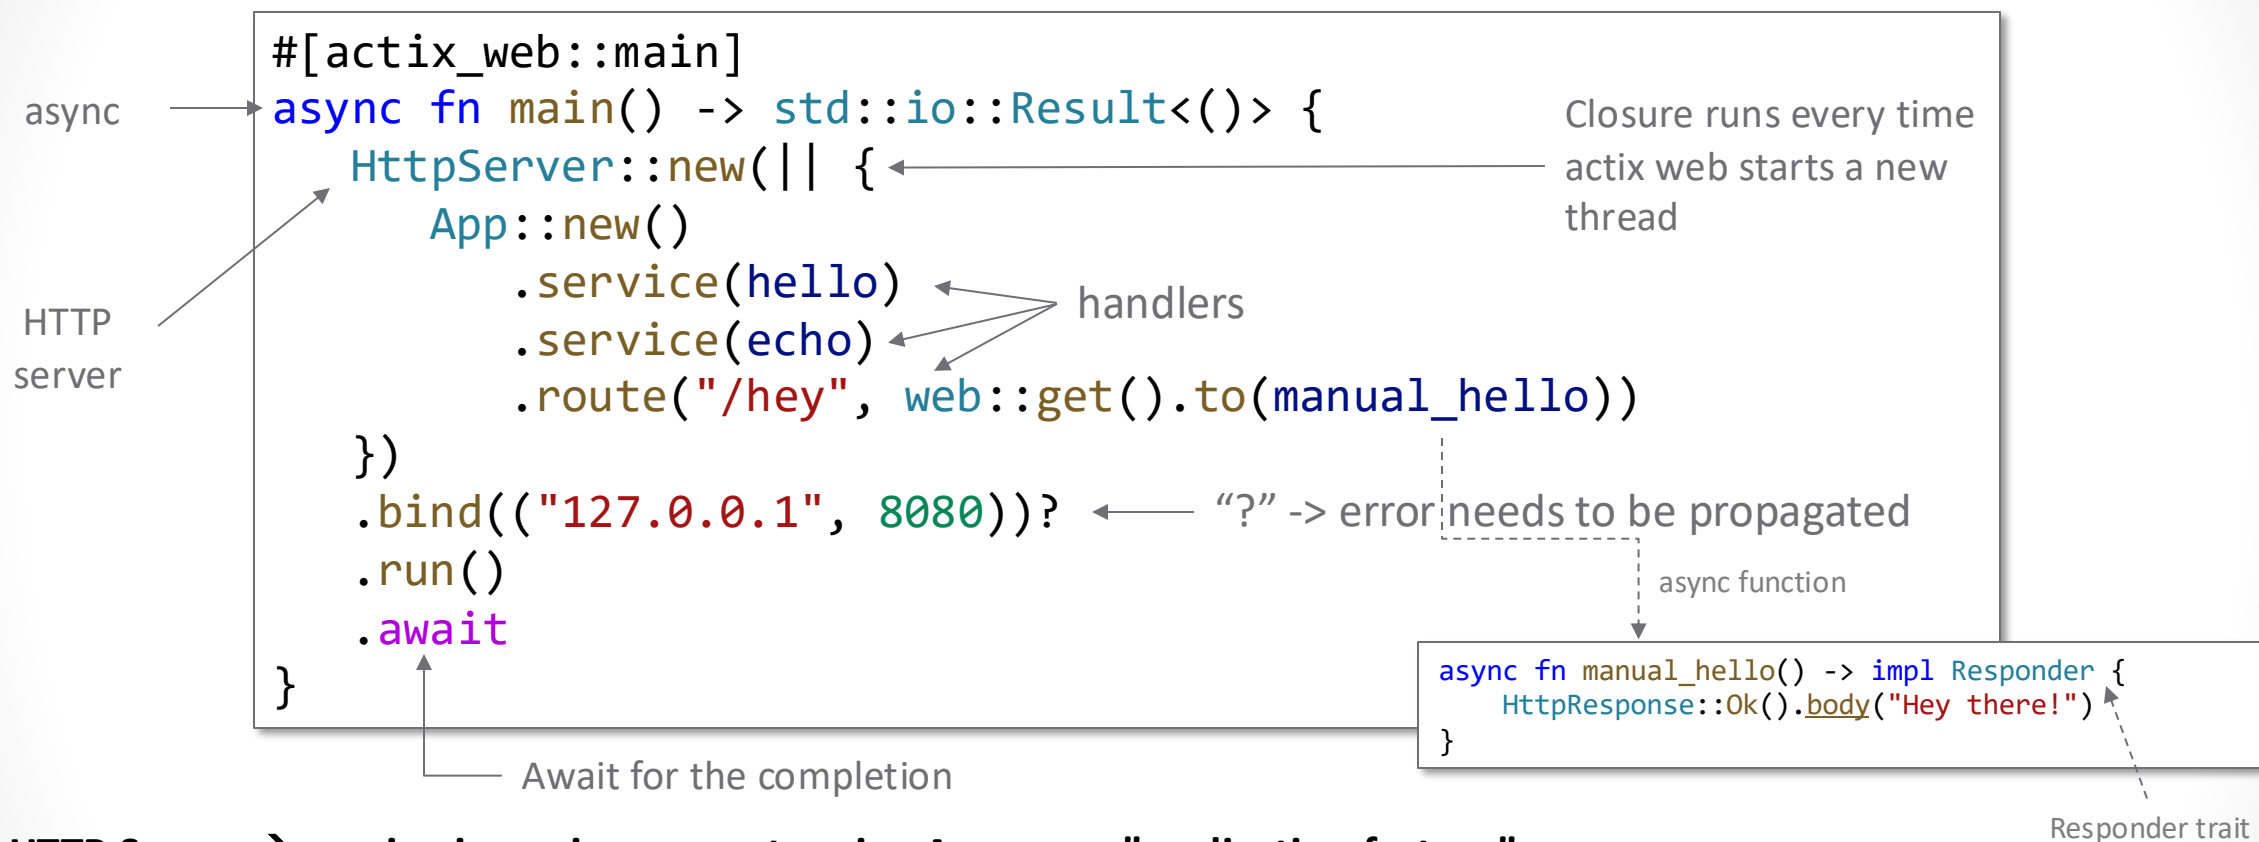
- Open-source
- Type safe
- Feature rich (HTTP, HTTP2, TLS, logging, ...)
- Powerful request routing
  - Can automatically extract data from incoming requests
- Support for multi-threading
- Middlewares (Logger, Session CORS)
- Developer friendly (Responders, Extractors, Form Handling, ...)
- ...

<https://github.com/actix/actix-web>

<https://actix.rs>, <https://actix.rs/docs/>

<https://www.techempower.com/benchmarks/>

# ACTIX WEB :: HELLO WORLD



**HTTP Server → serving incoming requests using App as an "application factory"**

**App ("application factory")**

- Registering routes for resources and middleware
- Storing application state across handlers with the same scope

**Cargo.toml, dependencies: actix-web = { version = "^4" }**

source: <https://actix.rs/docs/>



# ACTIX WEB :: HTTP SERVER

## HttpServer is Responsible for serving HTTP requests

- Accepts an application factory as a parameter
- It must be bound to a network socket with `HttpServer::bind()`
  - For example, with `.bind(("127.0.0.1", 8080))`?
  - This can fail if the socket is in used, therefore this error must be handled
- Multi-threaded
  - `.run()` returns a Server instance
    - It has to be **awaited** or spawned
  - It automatically starts a number of workers (default: the number of physical CPUs on the system)
    - Can be altered with `.workers(4)` method
  - Each worker threads processes its requests sequentially
    - A special care should be taken when blocking a thread for any reason
      - Use asynchronous functions and futures
    - async “handlers” are executed concurrently by worker threads
  - Each thread has a **separate instance of "App" to handle requests**
  - Concurrent access to shared data must be protected (e.g., with **Arc**)

# ACTIX WEB :: APPLICATION FACTORY

## "App"

- Registering routes for resources and middleware
- Stores application state shared across all handlers within the same scope
- Provides an application scope that acts as a namespace for all routes

```
use actix_web::{web, App, HttpServer, Responder};

async fn index() -> impl Responder {
    "Hello world!"
}

#[actix_web::main]
async fn main() -> std::io::Result<()> {
    HttpServer::new(|| {
        App::new().service(
            // prefixes all resources and routes attached to it...
            web::scope("/app")
            // ...so this handles requests for `GET /app/index.html`
            .route("/index.html", web::get().to(index)),
        )
    })
    .bind(("127.0.0.1", 8080))?
    .run()
    .await
}
```

Will match: /app, /app/

Will match: /app/index.html

# ACTIX WEB :: APP INSTANCE, AND SHARED STATE

## Shared State

- Shared with all routes and resource within the same scope
- Accessed with `web::Data<T>`
  - T is the type of the state

```
use actix_web::{get, web, App, HttpServer};

// This struct represents state
struct AppState {
    app_name: String,
}

#[get("/")]
async fn index(data: web::Data<AppState>) -> String {
    let app_name = &data.app_name; // <- get app_name
    format!("Hello {app_name}!")
}
// ...
```

(we still need the main)

```
// ...
#[actix_web::main]
async fn main() -> std::io::Result<()> {
    HttpServer::new(|| {
        App::new()
            .app_data(web::Data::new(AppState {
                app_name: String::from("Actix Web"),
            }))
            .service(index)
    })
    .bind(("127.0.0.1", 8080))?
    .run()
    .await
}
```

# ACTIX WEB :: APP INSTANCE AND MUTABLE STATE

## Mutable State

- `HttpServer` constructs an instance of `App` for each thread
- To share data between threads a shareable object should be used (e.g., `Send + Sync*`)
- Internally, `web::Data` uses `Arc` (Atomically Reference Counted\*\*)

```
use actix_web::{web, App, HttpServer};
use std::sync::Mutex;

struct MyAppData {
    // Mutex for modifyinh safely across threads
    request_count: Mutex<i32>,
}

async fn index(data: web::Data<MyAppData>) -> String {
    // get MutexGuard<'_, i32>
    let mut request_count = data.request_count.lock().unwrap();
    // now we have exclusive access by a thread
    *request_count += 1;

    // We need to return string
    format!("Request count: {total_request_count}")
}
```

\*see: <https://doc.rust-lang.org/nomicon/send-and-sync.html>

\*\*see: <https://doc.rust-lang.org/std/sync/struct.Arc.html>

```
#[actix_web::main]
async fn main() -> std::io::Result<()> {
    // Note: web::Data created before the closure below
    let data = web::Data::new(MyAppData {
        request_count: Mutex::new(0),
    });

    HttpServer::new(move || {
        // move the counter into the closure
        App::new()
            .app_data(data.clone()) // register the data
            .route("/", web::get().to(index))
    })
    .bind(("127.0.0.1", 8080))?
    .run()
    .await
}
```

# ACTIX WEB :: CONFIGURE

## Simplicity and reusability of App and web :: scope

```
use actix_web::{web, App, HttpResponse, HttpServer};

fn scoped_config(cfg: &mut web::ServiceConfig) {
    cfg.service(
        web::resource("/test")
            .route(web::get().to(|| async { HttpResponse::Ok().body("test") })))
            .route(web::head().to(HttpResponse::MethodNotAllowed)),
    );
}

fn config(cfg: &mut web::ServiceConfig) {
    cfg.service(
        web::resource("/app")
            .route(web::get().to(|| async { HttpResponse::Ok().body("app") })))
            .route(web::head().to(HttpResponse::MethodNotAllowed)),
    );
}
```

```
#[actix_web::main]
async fn main() -> std::io::Result<()> {
    HttpServer::new(|| {
        App::new()
            .configure(config)
            .service(web::scope("/api").configure(scoped_config))
    })
    .bind(("127.0.0.1", 8080))?
    .run()
    .await
}
```

# ACTIX WEB :: EXTRACTORS

## Request information access (extraction)

- Type safety
- Trait impl **FromRequest**
  - Many built-in extractor implementations  
([https://docs.rs/actix-web/latest/actix\\_web/trait.FromRequest.html#implementors](https://docs.rs/actix-web/latest/actix_web/trait.FromRequest.html#implementors))
- Can be accessed as an argument to a “handler” function

```
async fn index(path: web::Path<(String, String)>, json: web::Json<MyInfo>) -> impl Responder {  
    let path = path.into_inner();  
    format!("{}", path.0, path.1, json.id, json.username)  
}
```

## Examples

- Path, Query, Json (see next slides)
- Data (accessing pieces of application state)
- HttpRequest (parts of the request)
- String (conversion of a request payload to a String)

# ACTIX WEB :: EXTRACTORS :: PATH

## Path

- Information that is extracted from the request's path
- Parts of the path that are extractable are called “dynamic segments”, marked with { and }
- You can deserialize any variable segment from the path
- Usable also in combination with “**serde**” Deserialize trait

```
use actix_web::{get, web, App, HttpServer, Result};

/// extract path info from "/users/{user_id}/{friend}" url
/// {user_id} - deserializes to a u32
/// {friend} - deserializes to a String
#[get("/users/{user_id}/{friend}")] // <- define path parameters
async fn index(path: web::Path<(u32, String)>) -> Result<String> {
    let (user_id, friend) = path.into_inner();
    Ok(format!("Welcome {}, user_id {}", friend, user_id))
}

#[actix_web::main]
async fn main() -> std::io::Result<()> {
    HttpServer::new(|| App::new().service(index))
        .bind(("127.0.0.1", 8080))?
        .run()
        .await
}
```

# ACTIX WEB :: EXTRACTORS :: PATH

## Path

- Information that is extracted from the request's path (with `web::Path<Info>`)
- Parts of the path that are extractable are called “dynamic segments”, marked with { and }
- You can deserialize any variable segment from the path
- Usable also in combination with **“serde” Deserialize trait**

```
use actix_web::{get, web, Result};
use serde::Deserialize;

#[derive(Deserialize)]
struct Info {
    user_id: u32,
    friend: String,
}

/// extract path info using serde
#[get("/users/{user_id}/{friend}")] // <- define path parameters
async fn index(info: web::Path<Info>) -> Result<String> {
    Ok(format!("Welcome {}, user_id {}", info.friend, info.user_id ))
}

// .., main ...
```



# ACTIX WEB :: EXTRACTORS :: QUERY

## Query

- With `web::Query<T>`
- Used for request parameters
- Uses `serde_urlencoded` behind the scenes

```
use actix_web::{get, web, App, HttpServer};
use serde::Deserialize;

#[derive(Deserialize)]
struct Info {
    username: String,
}

// this handler gets called if the query deserializes into `Info` successfully
// otherwise a 400 Bad Request error response is returned
#[get("/")]
async fn index(info: web::Query<Info>) -> String {
    format!("Welcome {}", info.username)
}
```

# ACTIX WEB :: EXTRACTORS :: JSON

## Json

- With `web::Json<Info>`
- Deserialization of request body into a struct
- Note that type T must implement the Deserialize trait

```
use actix_web::{post, web, App, HttpServer, Result};
use serde::Deserialize;

#[derive(Deserialize)]
struct Info {
    username: String,
}

/// deserialize `Info` from request's body
#[post("/submit")]
async fn submit(info: web::Json<Info>) -> Result<String> {
    Ok(format!("Welcome {}!", info.username))
}
```

# ACTIX WEB :: EXTRACTORS :: DATA

## Data

- Access to application state
- With `web::Data<AppState>`
- Mutable access must be implemented

```
use actix_web::{post, web, App, HttpServer, Result};
use serde::Deserialize;

#[derive(Deserialize)]
struct Info {
    username: String,
}

/// deserialize `Info` from request's body
#[post("/submit")]
async fn submit(info: web::Json<Info>) -> Result<String> {
    Ok(format!("Welcome {}", info.username))
}
```

# ACTIX WEB :: HANDLERS :: CUSTOM TYPES

Returning custom types is possible with the **Responder** trait

```
use actix_web::{ body::BoxBody, http::header::ContentType, HttpRequest, HttpResponse, Responder, };
use serde::Serialize;

#[derive(Serialize)]
struct MyObj {
    name: &'static str,
}

// Responder
impl Responder for MyObj {
    type Body = BoxBody;

    fn respond_to(self, _req: &HttpRequest) -> HttpResponse<Self::Body> {
        let body = serde_json::to_string(&self).unwrap();

        // Create response and set content type
        HttpResponse::Ok()
            .content_type(ContentType::json())
            .body(body)
    }
}

async fn index() -> impl Responder {
    MyObj { name: "user" }
}
```

# ACTIX WEB :: HANDLERS :: DIFFERENT RESPONSES

**It is also possible to have a handler that returns different type of responses**

- “Either” type can be used

```
use actix_web::{Either, Error, HttpResponse};

type RegisterResult = Either<HttpResponse, Result<&'static str, Error>>;

async fn index() -> RegisterResult {
    if is_a_variant() {
        // choose Left variant
        Either::Left(HttpResponse::BadRequest().body("Bad data"))
    } else {
        // choose Right variant
        Either::Right(Ok("Hello!"))
    }
}
```

THE LAST SLIDE

**Questions?**